

Artificial neural networks - Exercise session 1

Supervised learning and generalization

2017-2018

1 The perceptron

The perceptron is the simplest one-layer network. It consists of R inputs connected to N neurons arranged in a single layer via weighted connections. These neurons have the `hardlim` transfer function, thus the output values are only 0 and 1, where the inputs can take on any value. The perceptron is used as a simple classification tool.

To create a perceptron we can use the command

```
net = newp(P,T,TF,LF)
```

where P and T are input and target vectors e.g. $P=[2 \ 1 \ -2 \ -1; \ 2 \ -2 \ 2 \ 1]$, $T=[0 \ 1 \ 0 \ 1]$. TF is the transfer function (typically 'hardlim'), LF represents the perceptron learning rule (for instance 'learnp'). In this case the number of neurons is set automatically.

The weights of the connections and the bias of the neurons are initially set to zero, which can be checked and changed with the commands

<code>net.IW{1,1}</code>	Returns the weights of the neuron(s) in the first layer
<code>net.b{1,1}</code>	Returns the bias of the neuron(s) in the first layer
<code>net.IW{1,1} = rand(1,2);</code>	Assigns random weights in [0,1]
<code>net.b{1,1} = rands(1);</code>	Assigns random bias in [-1,1]

One can initialize the network with a single command by issuing:

```
net = init(net);
```

all weights and biases of the perceptron will be initialized to zero.

We can teach a perceptron to perform certain tasks which are defined by pairs of inputs and outputs. A perceptron can learn only linearly separable tasks: inputs belonging to different classes are separated by a hyperplane in the input space (decision boundary). In two dimensions the decision boundary is a line.

We can use the function `train` to let the perceptron perform the classification task. In this case the learning occurs in batch mode:

```
[net,tr_descr] = train(net,P,T);
```

here the second argument is a description of the learning process.

To set the number of iterations or epochs, we can use the command

```
net.trainParam.epochs = 20;
```

Finally, after training we can simulate the network on new data with the function `sim`:

```
sim(net,Pnew)
```

with P_{new} an input vector. For our example P_{new} has to be a (column) vector of length 2, with elements between -2 and 2, e.g. $P_{new} = [1; -0.3]$. Multiple input vectors can be fed at once to the network by putting them together in an array.

Demos

The following demos can be run from the MATLAB prompt. For the graphical demos follow the instructions on screen. For non-graphical demos you can run them interactively with the command `playshow name_of_demo`, and see and edit (a copy of) the code with `edit name_of_demo`

<code>nnd4db</code>	decision boundary (2d input)
<code>nnd4pr</code>	perceptron learning rule (2d input)
<code>playshow demop1</code>	classification with 2d input perceptron
<code>playshow demop4</code>	classification with outlier (2d input)
<code>playshow demop5</code>	classification with outlier using normalized perceptron learning rule (2d input)
<code>playshow demop6</code>	linearly non-separable input vectors (2d input)

Exercises

Create a perceptron and train it with examples (see demos). Can you always train it perfectly?

Functions and commands

<code>newp(P, T, TF, LF)</code>	Creates a perceptron, based on input values P, target vector T, and with transfer function TF and learning function LF.
<code>init(net)</code>	Initializes the weights and biases of the perceptron.
<code>adapt(net, P, T)</code>	Trains the network using inputs P, targets T and some online learning algorithm.
<code>train(net, P, T)</code>	Trains the network using inputs P, targets T and some batch learning algorithm.
<code>sim(net, Ptest)</code>	Simulates the perceptron using inputs Ptest.
<code>learnp, learnpn</code>	Perceptron and normalized perceptron learning rules
<code>hardlim</code>	Transfer function

2 Backpropagation in feedforward multi-layer networks

A general feedforward network consists of at least one layer, and it can also contain an arbitrary number of hidden layers. Neurons in a given layer can be defined by any transfer function. In the hidden layers usually nonlinear functions are used, e.g. `tansig` or `logsig`, and in the output layer `purelin`. The only important condition is that there is no feedback in the network, neither delay.

In MATLAB one can create such a network object by e.g. the following command:

```
net = feedforwardnet(numN, trainAlg);
```

This will create a network of one hidden layer with corresponding `numN` neurons, which will use the `trainAlg` algorithm for training (e.g. `traingd`). This network can be trained using the `train` function:

```
net = train(net, P, T);
```

Finally the network can be simulated in two ways:

```
sim(net, P);
```

or,

```
Y = net(P);
```

Some available training algorithms are:

<code>traingd</code>	gradient descent
<code>traingda</code>	gradient descent with adaptive learning rate
<code>traincgf</code>	Fletcher-Reeves conjugate gradient algorithm
<code>traincgp</code>	Polak-Ribiere conjugate gradient algorithm
<code>trainbfg</code>	BFGS quasi Newton algorithm (quasi Newton)
<code>trainlm</code>	Levenberg-Marquardt algorithm (adaptive mixture of quasi Newton and steepest descent algorithms)

To analyze the efficiency of training one can use the function `postreg` (can be found in Toledo) which calculates and visualizes regression between targets and outputs. For a network `net` trained with a sequence of examples `P` and targets `T` we have:

```
a=sim(net,P);
[m,b,r]=postreg(a,T);
```

where `m` and `b` are the slope and the y-intercept of the best linear regression respectively. `r` is a correlation between targets `T` and outputs `a`.

Demos

<code>nnd11nf</code>	network function
<code>nnd11bc</code>	backpropagation calculation
<code>nnd11fa</code>	function approximation
<code>nnd12sd1</code>	steepest descent backpropagation
<code>nnd12sd2</code>	steepest descent backpropagation with various learning rates
<code>nnd12mo</code>	steepest descent with momentum
<code>nnd12v1</code>	steepest descent with variable learning rate
<code>nnd12cg</code>	conjugate gradient backpropagation
<code>nnd9mc</code>	comparison between steepest descent and conjugate gradient

Exercises

- Function approximation: comparison of various algorithms:
Take $y = \sin(x^2)$ for $x=0:0.05:3\pi$ as a simple nonlinear function and try to approximate it using a neural network with one hidden layer. Use different algorithms. How does gradient descent perform compared to other training algorithms? Use following examples as a basis:
`algorlml` (can be found in Toledo) Script that compares the performance of Levenberg-Marquardt 'trainlm' and quasi-Newton backpropagation 'trainbfg' algorithms
- Learning from noisy data: generalization
The same as in the previous exercise, but now add noise to the data (small random numbers to each datapoint). Compare the performance of the network with noiseless data. You may have to increase the number of data.

3 Understanding Bayesian Inference: the Case of Network Weights

For simplicity of exposition, we begin by considering the training of a network for which the architecture is fixed in advance. More precisely we focus on the case of a one-neuron network. In the absence of any data, the distribution over weight values is described by a prior distribution denoted $p(w)$, where w is the vector of adaptive weights (normally also biases, but in our example we will consider the case without bias). We also denote by D the available dataset. Once we observe the data, we can write the expression for the posterior probability distribution of the weights using Bayes' theorem:

$$p(w|D) = \frac{p(D|w)p(w)}{p(D)} \quad (1)$$

where $p(D)$ is a normalization factor ensuring that $p(w|D)$ gives unity when integrated over the whole weight space. $p(D|w)$ corresponds to the likelihood function used in maximum likelihood techniques.

Since in the beginning we do not know anything about the data, the prior distribution of weights is set to (for example) a Gaussian distribution. The expression of the prior is then:

$$p(w) = \left(\frac{\alpha}{2\pi}\right)^{\frac{w}{2}} \exp\left(-\frac{\alpha}{2}||w||^2\right) \quad (2)$$

with W the number of weights, and α is the inverse of the variance. For simplicity we will choose $\alpha = 1$.

This choice of the prior distribution encourages weights to be small rather than large, which is a requirement for achieving smooth network mappings. So when $\|w\|$ is large, the parameter of the exponential is large, and thus $p(w)$ is small (small probability that this is the correct choice of weights. Things are reversed for small $\|w\|$).

A concrete example: binary classification We consider the case where input vectors are 2-dimensional $x = (x_1, x_2)$, and we have four data points in our dataset as in figure 1. We take a network of a single neuron (compare to the perceptron), so there is only a single layer of weights, and choose the logistic function as transfer function:

$$y(x; w) = \frac{1}{1 + \exp(-w^T x)} \quad (3)$$

The weight vector $w = (w_1, w_2)$ is two-dimensional and there is no bias parameter. We choose a Gaussian prior distribution for the weights (with $\alpha = 1$). This prior distribution is plotted in figure 2.

The data points can belong to one of the two classes (cross or circle), the output y giving the membership to one of these classes. The likelihood function $p(D|w)$ in Bayes' theorem will be given by a product of factors, one for each data point, where each factor is either y or $(1 - y)$ according to whether the data point belongs to the first or the second class.

First we consider just the points labeled (1) and (2). Then we consider all four points and recompute the posterior distribution of the weights.

Note: For an example of this case, run demo `bayesNN`, that you can find in the Exercise session 1 folder on Toledo → Assignments .

After training with only the first two data points, we see that the network function is a sigmoidal ridge (w_1 and w_2 control the orientation and the slope of this sigmoid). Weight vectors from approximately half of the weight space will have probabilities close to zero, as they represent decision boundaries with wrong orientation. When using all 4 data points, there is no boundary to classify all 4 points correctly. The most probable solution is the one corresponding to the peak point of the sigmoid, the others having quite low probabilities (so the posterior distribution of the weights is relatively narrow).

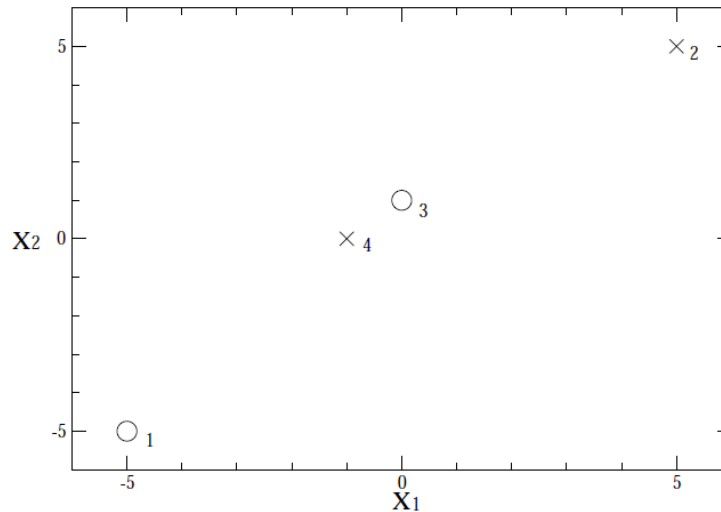


Figure 1: The four numbered data points in the dataset.

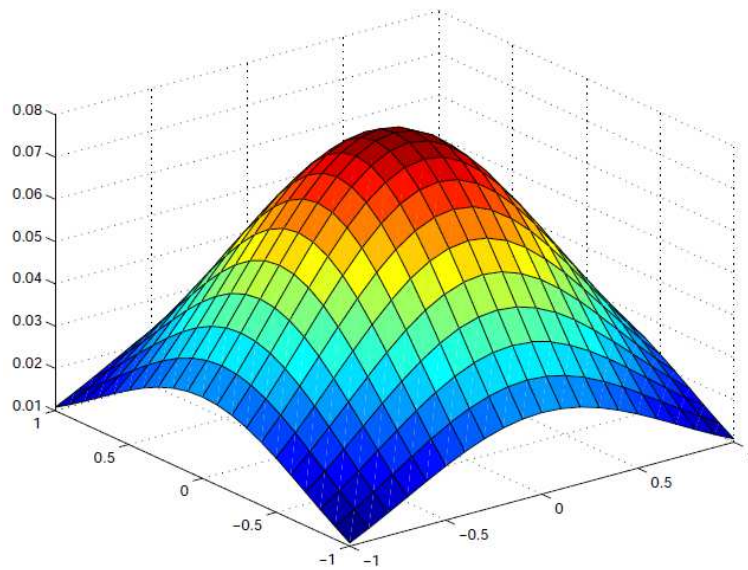


Figure 2: The prior distribution: a Gaussian.

4 Bayesian Inference of Network? hyperparameters

From an optimization point of view the training is performed by iteratively adjusting w so as to minimize an objective function $M(w)$. In a maximum likelihood framework, M is taken to be the error function $E_D(w)$ that depends on the data D . A common prescription to avoid overfitting is to include in M a regularization term $E_W(w)$ (a.k.a. weight decay) to favor small values of the parameter vectors, as discussed above. In particular, if we let

$$M(w) = \beta E_D(w) + \alpha E_W(w) \quad (4)$$

we can immediately give to this a Bayesian interpretation. In fact it is not difficult to show that βE_D can be understood as minus the log likelihood for a noise model whereas αE_W can be understood as minus the log prior on the parameter vector. Correspondingly, the process of finding the optimal weight vector minimizing M can be interpreted as a Maximum a Posteriori (MAP) estimation. Notice that we have implicitly considered α and β fixed here. However these are hyperparameters that control the complexity of the model. Once more within a Bayesian framework one can apply the rules of probability and find these parameters accordingly. In MATLAB for a general feedforward neural network this can be accomplished by means of `trainbr`. This training function updates the weight according to Levenberg-Marquardt optimization. It minimizes (4) with respect to w and determines at the same time the correct combination of the two terms so as to produce a network that generalizes well.

Exercises

- Use `trainbr` to analyse the artificial datasets of section 2 and compare it with the training functions seen previously. Compare the test errors. Consider overparametrized networks (many neurons): do you see any improvement with `trainbr`?

5 Report

Based on the previous exercises of section 2 and 4, write a report of maximum 2 pages (including text and figures) to discuss speed, overfitting and generalization of different learning schemes.

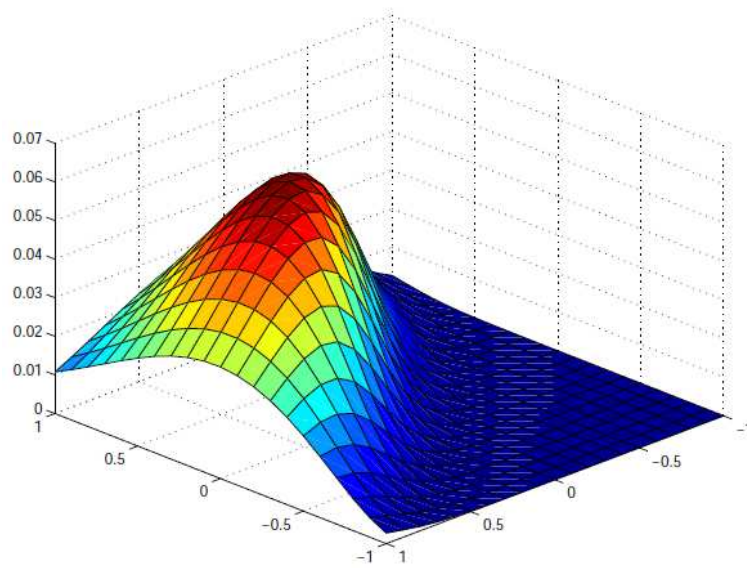


Figure 3: The distribution after presenting all four data points. Half of the weight space has become very improbable.

References

- [1] H. Demuth and M. Beale, Neural Network Toolbox (user's guide),
<http://www.mathworks.com/access/helpdesk/help/toolbox/nnet/nnet.shtml>
- [2] C. M. Bishop, Neural Networks for Pattern Recognition, Oxford University Press.
- [3] Lecture slides and references therein.